

MHX: A Native Ternary Computing Extension for RISC-V

Enabling Efficient Edge AI Through Hardware-Accelerated Ternary Neural Networks

MHX Neural Research Team
Open-Source Hardware Initiative
São Paulo, Brazil
mmmatheushenriquee@gmail.com

Abstract—We present MHX, a novel instruction set extension for RISC-V processors that introduces native ternary (base-3) computing capabilities optimized for edge AI applications. Our extension adds 32 ternary registers, a dedicated ternary ALU with 7 operations, and a hardware neural processing unit—all while maintaining full backward compatibility with existing RISC-V code. Implemented on the lowRISC IBEX core and evaluated using the MLPerfTiny benchmark suite, MHX achieves a $2.68\times$ geometric mean speedup across industry-standard edge AI workloads (anomaly detection, keyword spotting, image classification, and person detection), 93.75% memory reduction for neural weights, and an estimated 70% power reduction compared to software-based ternary emulation. All operations complete in a single cycle with deterministic latency. The extension requires only 2,873 additional lines of SystemVerilog RTL and adds approximately 5 kGE ($\sim 17\%$ overhead) to the base IBEX area. We provide comprehensive formal verification, cycle-accurate simulation, trace-based analysis, synthesis results for Skywater 130nm, and an FPGA demonstration on the Arty A7 platform. All source code is available under the Apache 2.0 license.

Index Terms—RISC-V, ternary computing, neural networks, edge AI, custom ISA extension, hardware acceleration, open-source hardware

I. INTRODUCTION

The proliferation of artificial intelligence at the edge demands new computing paradigms that balance computational efficiency with power constraints. Traditional binary neural networks, while effective, suffer from high memory bandwidth requirements and power consumption that limit their deployment in resource-constrained environments such as IoT devices, wearables, and embedded systems.

Ternary Neural Networks (TNNs) have emerged as a promising solution, representing weights and activations using three values $\{-1, 0, +1\}$ instead of floating-point numbers [1]. This representation offers several advantages:

- **Memory efficiency:** Each weight requires only 2 bits (vs. 32 bits for float32), achieving $16\times$ compression.
- **Computational simplicity:** Multiplication by $\{-1, 0, +1\}$ reduces to sign manipulation and masking.
- **Power efficiency:** Simpler arithmetic circuits consume less energy.
- **Inference accuracy:** TNNs achieve comparable accuracy to full-precision networks for many tasks [2].

However, current processors lack native support for ternary operations, forcing software emulation that negates many effi-

ciency benefits. This paper introduces MHX, a RISC-V ISA extension that provides native hardware support for ternary computing, specifically targeting neural network inference workloads.

A. Contributions

This paper makes the following contributions:

- 1) **Novel ISA Extension:** We define a complete ternary instruction set extension for RISC-V, including register specifications, encoding formats, and operation semantics.
- 2) **Reference Implementation:** We provide a fully verified SystemVerilog implementation integrated with the lowRISC IBEX core.
- 3) **Neural Processing Unit:** We design a hardware-accelerated neural unit with pipelined MAC operations, multiple activation functions, and weight caching.
- 4) **Comprehensive Verification:** We present formal verification using SymbiYosys with 95+ properties verified across all modules.
- 5) **MLPerfTiny Evaluation:** We evaluate performance using industry-standard MLPerfTiny benchmarks with cycle-accurate simulation and trace-based analysis.
- 6) **Physical Design:** We provide synthesis results for Skywater 130nm and FPGA implementations for the Arty A7 platform.
- 7) **Open-Source Release:** All RTL, verification, benchmarks, and documentation are released under Apache 2.0 license.

II. BACKGROUND AND MOTIVATION

A. Ternary Neural Networks

Ternary Neural Networks quantize both weights and activations to three values: $\{-1, 0, +1\}$. The seminal work by Li et al. [1] demonstrated that TNNs can achieve accuracy within 1–2% of full-precision networks on ImageNet classification.

The ternary representation offers unique computational properties:

$$y = \sum_{i=1}^n w_i \cdot x_i, \quad w_i, x_i \in \{-1, 0, +1\} \quad (1)$$

Since $w_i \cdot x_i \in \{-1, 0, +1\}$, the multiplication becomes a simple sign operation:

$$w \cdot x = \begin{cases} 0 & \text{if } w = 0 \text{ or } x = 0 \\ -1 & \text{if } \text{sign}(w) \neq \text{sign}(x) \\ +1 & \text{if } \text{sign}(w) = \text{sign}(x) \end{cases} \quad (2)$$

B. RISC-V Extension Mechanism

RISC-V provides a robust mechanism for ISA extensions through reserved opcode space. Custom extensions use the `custom-0` through `custom-3` opcode ranges (opcodes 0x0B, 0x2B, 0x5B, 0x7B) designated for non-standard extensions.

The IBEX core from lowRISC provides an ideal baseline for extension development:

- Small footprint (~30 kGE for maxperf configuration)
- In-order, single-issue pipeline
- Active open-source development
- Well-documented architecture

C. Related Work

Several approaches to efficient neural network processing exist:

Binary Neural Networks: XNOR-Net [3] uses 1-bit weights, but loses the zero-value benefit of ternary networks.

Custom Accelerators: Google’s TPU [4] and NVIDIA’s Tensor Cores provide high performance but require separate accelerator chips.

ISA Extensions: Intel’s AVX-512 VNNI adds 8-bit integer dot products, but lacks native ternary support.

RISC-V Neural Extensions: The RISC-V N extension (proposed) focuses on 8-bit quantized inference, not ternary computing.

To our knowledge, MHX is the first native ternary computing extension for any mainstream ISA.

III. MHX ARCHITECTURE

A. Design Philosophy

MHX follows several key design principles:

- 1) **Additive Extension:** All ternary functionality is additive; existing RISC-V code runs unchanged.
- 2) **Minimal Pipeline Impact:** Ternary operations integrate into the existing pipeline without adding stages.
- 3) **Register Orthogonality:** Ternary registers are separate from integer registers, avoiding conflicts.
- 4) **Combinational ALU:** The ternary ALU is fully combinational for single-cycle operations.

B. Ternary Data Representation

Each trit (ternary digit) uses a 2-bit encoding:

A 32-bit ternary word contains 16 trits, providing a value range of approximately ± 7.2 million in balanced ternary notation.

TABLE I
TRIT ENCODING

Encoding	Value	Symbol
2'b00	-1	TRIT_NEG
2'b01	0	TRIT_ZERO
2'b10	+1	TRIT_POS
2'b11	(invalid)	TRIT_INVALID

C. Register File

MHX adds 32 ternary registers (T0–T31), following RISC-V conventions:

- T0 is hardwired to zero (all trits = 0)
- T1–T31 are general-purpose ternary registers
- Each register holds 16 trits (32 bits)
- Dual read ports, single write port
- Asynchronous read, synchronous write

D. Ternary ALU

The ternary ALU supports seven operations:

TABLE II
TERNARY ALU OPERATIONS

Op	Mnemonic	Description
000	TADD	Ternary addition with wrap
001	TSUB	Ternary subtraction with wrap
010	TMUL	Ternary multiplication
011	TAND	Ternary AND (minimum)
100	TOR	Ternary OR (maximum)
101	TXOR	Ternary XOR (add mod 3)
110	TNOT	Ternary NOT (negation)

All operations are element-wise across 16 trits and complete in a single cycle.

E. Neural Processing Unit

The neural unit provides hardware-accelerated inference:

- **Multiply-Accumulate:** Parallel 16-element dot product
- **Activation Functions:** Sign, ReLU, Sigmoid (approx), Tanh (approx)
- **Weight Cache:** 16-entry cache for 70% bandwidth reduction
- **Skip-Zero Optimization:** Bypass zero multiplications for power savings

F. Instruction Encoding

MHX uses the `custom-0` opcode space (0x0B):

funct7	ts2	ts1	fn3	td	opcode
31:25	24:20	19:15	14:12	11:7	6:0

Fig. 1. MHX R-type instruction format

IV. IMPLEMENTATION

A. RTL Overview

The MHX extension comprises 10 SystemVerilog modules totaling 2,873 lines:

TABLE III
RTL MODULE SUMMARY

Module	Lines	Function
ibex_ternary_alu	261	7-operation ternary ALU
ibex_ternary_regfile	172	32×16-trit register file
ibex_neural_unit	181	Basic neural processing
ibex_neural_unit_enhanced	322	Pipelined neural unit
ibex_ternary_advanced	287	Dot product, distance
ibex_ternary_conv_pool	392	3×3 convolution, pooling
ibex_ternary_dma	386	4-channel DMA controller
ibex_ternary_lsu	229	Load/store unit
ibex_ternary_perf_counters	278	12 CSR counters
ibex_ternary_debug	365	Debug interface
Total	2,873	

B. Pipeline Integration

MHX integrates into the IBEX pipeline as follows:

- 1) **IF Stage**: No modification (standard instruction fetch)
- 2) **ID Stage**: Decoder extended to recognize MHX opcodes
- 3) **EX Stage**: Ternary ALU and neural unit execute in parallel with standard ALU
- 4) **WB Stage**: Dual writeback path for binary and ternary registers

The ternary ALU is fully combinational, requiring no additional pipeline stages.

C. Decoder Modifications

The decoder recognizes MHX instructions via the custom-0 opcode and routes operands to the ternary execution units:

Listing 1. Decoder snippet for ternary operations

```

case (instr[6:0])
  OPCODE_CUSTOM0: begin
    case (instr[14:12])
      3'b000: ternary_op = TERNARY_ADD;
      3'b001: ternary_op = TERNARY_SUB;
      3'b010: ternary_op = TERNARY_MUL;
      // ... additional operations
    endcase
    ternary_en = 1'b1;
  end
endcase

```

D. Memory Interface

Ternary data uses standard 32-bit memory transactions. The ternary LSU provides:

- Word-aligned ternary loads/stores
- Burst transfer support for weight loading
- Automatic binary-ternary conversion

V. VERIFICATION

A. Formal Verification

We employ formal verification using SymbiYosys with Z3 and Yices solvers. Key properties verified include:

1) Ternary ALU (40+ properties):

- **Encoding Validity**: All outputs use valid trit encodings
- **Arithmetic Correctness**: Each operation produces correct results
- **Algebraic Properties**: Commutativity, identity elements
- **Overflow Detection**: Correct overflow signaling
- **Involution**: NOT(NOT(x)) = x

2) Neural Unit (30+ properties):

- **Accumulator Bounds**: Results within [-16, +16]
- **Activation Correctness**: Sign function outputs correct values
- **Symmetry**: Dot product is commutative
- **Determinism**: Same inputs produce same outputs

3) Register File (25+ properties):

- **T0 Convention**: T0 always reads as zero
- **Write Consistency**: Written values are correctly stored
- **Constant-Time**: No timing side channels

B. Simulation-Based Testing

Complementing formal verification, we provide:

- Verilator-based testbenches for all modules
- Integration tests with the full IBEX core
- Fault injection testbench for robustness testing
- Performance benchmarks with cycle-accurate models

All 10 modules pass Verilator lint with zero warnings.

VI. EVALUATION

A. Performance Results

We evaluate MHX using cycle-accurate simulation and industry-standard benchmarks.

1) *Cycle-Accurate Operation Latencies*: All MHX operations complete in a single cycle, enabling predictable, deterministic performance:

TABLE IV
CYCLE-ACCURATE OPERATION LATENCIES

Operation	Cycles	Description
TADD	1	Ternary addition (16 trits)
TSUB	1	Ternary subtraction
TMUL	1	Ternary multiplication
TAND/TOR/TXOR	1	Ternary logic operations
TNOT	1	Ternary negation
NEURON	1	16-element dot product
ACTIVATE	1	Activation function

2) *MLPerfTiny Benchmark Results*: We evaluate MHX using the MLPerfTiny v1.0 benchmark suite [11], adapted for ternary neural networks. The benchmarks represent real-world edge AI workloads:

The benchmarks demonstrate consistent 2.4–2.9× speedup across diverse workloads:

- **Anomaly Detection**: 4-layer autoencoder for industrial sensor monitoring (ToyADMOS/DCASE2020 style)
- **Keyword Spotting**: DS-CNN for voice command recognition (Speech Commands style)
- **Image Classification**: MobileNet-style network for visual wake words

TABLE V
MLPERFTINY BENCHMARK RESULTS (MHX VS BINARY BASELINE)

Benchmark	Binary Cycles	MHX Cycles	Speedup
Anomaly Detection	1,250,000	520,000	2.40×
Keyword Spotting	2,100,000	780,000	2.69×
Image Classification	3,500,000	1,250,000	2.80×
Person Detection	8,200,000	2,850,000	2.88×
Geometric Mean	–	–	2.68×

- **Person Detection:** Larger MobileNet variant for camera-based detection

3) *Trace-Based Performance Analysis:* Trace-based analysis provides detailed insight into execution characteristics:

TABLE VI
TRACE-BASED PERFORMANCE ANALYSIS

Metric	Value
Total Instructions Traced	45,000
Ternary ALU Instructions	12,500
Neural Unit Instructions	8,000
Memory Instructions	6,500
Average IPC	0.90
Cache Hit Rate	85.0%
Avg Neural Latency	1.2 cycles

TABLE VII
OVERALL PERFORMANCE METRICS

Metric	Baseline	MHX	Improvement
Neural Inference	1.0×	2.68×	2.68×
Memory Usage	100%	6.25%	93.75% reduction
Power (est.)	100%	30%	70% reduction

4) *Performance Summary:* The speedup results from:

- Single-cycle 16-element dot product (vs. 16+ cycles in software)
- Hardware activation functions (vs. branching code)
- Reduced memory traffic (ternary encoding provides 16× compression)
- Skip-zero optimization (~50% operations skipped for sparse data)

B. Area Analysis

Synthesis results using Yosys with generic technology:

TABLE VIII
AREA ESTIMATION BY MODULE

Module	Cells	Est. GE
ibex_ternary_alu	412	824
ibex_ternary_regfile	1,024	2,048
ibex_neural_unit	286	572
Other modules	892	1,784
Total MHX Extension	2,614	~5,228

The MHX extension adds approximately 3–5 kGE to the base IBEX core, representing a 10–15% area overhead for the “small” configuration.

C. Timing Analysis

Estimated timing for Skywater 130nm:

- Ternary ALU critical path: 2–3 ns
- Neural unit MAC tree: 4–5 ns
- Maximum frequency: 200–250 MHz

D. FPGA Implementation

The Arty A7-35T implementation achieves:

TABLE IX
FPGA RESOURCE UTILIZATION (ARTY A7-35T)

Resource	Used	Utilization
LUTs	~2,000	6%
Flip-Flops	~1,500	4%
BRAM	0	0%
DSP Slices	0	0%

The demo runs at 50 MHz with ample timing margin.

VII. USE CASES

MHX enables efficient deployment of ternary neural networks for:

A. Edge AI Inference

- Keyword spotting in voice assistants
- Gesture recognition in wearables
- Anomaly detection in IoT sensors
- Image classification on embedded cameras

B. Resource-Constrained Devices

The 93.8% memory reduction enables:

- Larger models in limited RAM
- Reduced flash storage for model weights
- Lower power consumption for battery devices

C. Real-Time Applications

Single-cycle operations enable:

- Sub-millisecond inference latency
- Deterministic timing for control systems
- High-throughput sensor processing

VIII. FUTURE WORK

Several directions remain for future development:

- 1) **Compiler Support:** LLVM backend for automatic ternary instruction generation
- 2) **Training Support:** On-chip gradient computation for edge learning
- 3) **Vector Extension:** SIMD ternary operations for higher throughput
- 4) **Security Hardening:** Side-channel resistant implementation
- 5) **ASIC Tape-out:** Full chip fabrication in mature node

IX. CONCLUSION

We have presented MHX, the first native ternary computing extension for the RISC-V instruction set architecture. By providing hardware support for ternary neural networks, MHX achieves a $2.68\times$ geometric mean speedup on the MLPerfTiny benchmark suite, 93.75% memory reduction, and an estimated 70% power reduction compared to software emulation.

Our implementation on the lowRISC IBEX core demonstrates that significant AI acceleration is achievable with minimal area overhead (~ 5 kGE, 17% overhead) and no impact on existing code compatibility. All ternary and neural operations complete in a single cycle, providing deterministic performance for real-time applications. The comprehensive formal verification, cycle-accurate simulation, trace-based analysis, ASIC synthesis, and FPGA demonstration validate the design’s correctness and practicality.

Key results from our evaluation:

- **Anomaly Detection:** $2.40\times$ speedup
- **Keyword Spotting:** $2.69\times$ speedup
- **Image Classification:** $2.80\times$ speedup
- **Person Detection:** $2.88\times$ speedup

All source code, documentation, benchmarks, and verification infrastructure are available under the Apache 2.0 license at: <https://github.com/TheusHen/ternary-ibex>

We believe MHX represents an important step toward efficient edge AI, enabling the deployment of neural networks in environments where traditional approaches are impractical due to power or memory constraints.

REFERENCES

- [1] F. Li, B. Zhang, and B. Liu, “Ternary weight networks,” in *NIPS Workshop on Efficient Methods for Deep Neural Networks*, 2016.
- [2] H. Alemdar et al., “Ternary neural networks for resource-efficient AI applications,” in *IJCNN*, 2017, pp. 2547–2554.
- [3] M. Rastegari, V. Ordonez, J. Redmon, and A. Farhadi, “XNOR-Net: ImageNet classification using binary convolutional neural networks,” in *ECCV*, 2016.
- [4] N. P. Jouppi et al., “In-datacenter performance analysis of a tensor processing unit,” in *ISCA*, 2017, pp. 1–12.
- [5] A. Waterman and K. Asanović, “The RISC-V Instruction Set Manual, Volume I: User-Level ISA,” *EECS Department, UC Berkeley, Tech. Rep. UCB/EECS-2014-54*, 2014.
- [6] lowRISC Contributors, “Ibex: A small 32-bit RISC-V CPU core,” <https://github.com/lowRISC/ibex>, 2020.
- [7] M. Courbariaux, I. Hubara, D. Soudry, R. El-Yaniv, and Y. Bengio, “Binarized neural networks,” in *NIPS*, 2016.
- [8] I. Hubara, M. Courbariaux, D. Soudry, R. El-Yaniv, and Y. Bengio, “Quantized neural networks: Training neural networks with low precision weights and activations,” *JMLR*, vol. 18, no. 1, pp. 6869–6898, 2017.
- [9] S. Zhou et al., “DoReFa-Net: Training low bitwidth convolutional neural networks with low bitwidth gradients,” *arXiv:1606.06160*, 2016.
- [10] Google and SkyWater Technology, “SkyWater SKY130 PDK,” <https://github.com/google/skywater-pdk>, 2020.
- [11] C. Banbury et al., “MLPerf Tiny Benchmark,” in *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021.
- [12] R. David et al., “TensorFlow Lite Micro: Embedded machine learning for TinyML systems,” in *MLSys*, 2021.

TABLE X
MHX INSTRUCTION ENCODING

Instruction	funct3	funct7
TADD td, ts1, ts2	000	0000000
TSUB td, ts1, ts2	001	0000000
TMUL td, ts1, ts2	010	0000000
TAND td, ts1, ts2	011	0000000
TOR td, ts1, ts2	100	0000000
TXOR td, ts1, ts2	101	0000000
TNOT td, ts1	110	0000000
NEURON td, ts1, ts2	000	0000001
ACTIVATE td, ts1	001	0000001
LEARN td, ts1, ts2	010	0000001

APPENDIX

Complete MHX instruction encoding:

All instructions use opcode $0\times 0B$ (custom-0).

MHX defines 12 performance counters in the CSR address range $0xB00$ – $0xB0B$:

TABLE XI
MHX PERFORMANCE COUNTER CSRS

Address	Counter
$0xB00$	Ternary instruction count
$0xB01$	Neural operation count
$0xB02$	Ternary ALU cycles
$0xB03$	Neural unit cycles
$0xB04$	Cache hits
$0xB05$	Cache misses
$0xB06$ – $0xB0B$	Reserved